

# Chapter 6 Problem 9

Andira Putri

In this exercise, we will predict the number of applications received using the other variables in the College data set.

a.) Split the data into a training set and a test set.

```
library(ISLR)
data(College) #contains 777 obs
set.seed(1)
train=sample(777,600) #600 obs for training
train.set=College[train,]
test.set=College[-train,] #177 obs for testing
```

b.) Fit a linear model using least squares on the training set, and report the test error obtained.

- Background on Least Squares Regression

Least squares regression attempts to find coefficients that minimize residual sum of squares (RSS). Ol' reliable.

```
#fit least squares model
lin.reg=lm(Apps~.,data=train.set)
#use test set on least squares model, make predictions
lin.pred=predict(lin.reg,test.set)
#calculate test error, just the general MSE formula
MSE.lin=mean((test.set$Apps-lin.pred)^2)
MSE.lin #And our test error is...
```

```
## [1] 1101792
```

c.) Fit a ridge regression model on the training set, with the tuning parameter chosen by cross-validation, and report the test error obtained.

- Background on Ridge Regression:

Ridge regression works almost like least squares regression, but the coefficients are estimated by minimizing a quantity slightly different than least squares. Least squares attempts to minimize just RSS. Ridge regression, on the contrary, attempts to minimize:

$$RSS + \lambda \sum_{i=1}^p \beta_i^2$$

The term  $\lambda \sum_{i=1}^p \beta_i^2$  is called a shrinking penalty, and it differentiates ridge regression from least squares.  $\lambda \geq 0$  is a tuning parameter that controls the impact of the shrinking penalty on the regression coefficient estimates. When  $\lambda = 0$ , the penalty has no effect, and it's just a least squares problem. As  $\lambda \rightarrow \infty$ , the ridge regression estimates will shrink towards (but never reaches) zero since the impact of the penalty is so large. Unlike least squares, ridge regression generates several sets of coefficients from varying values of  $\lambda$ . Selecting a good value for  $\lambda$  is important, and we do this via cross-validation methods.

Noteworthy:  $\|\beta_2\| = \sqrt{\sum_{i=1}^j \beta_i^2}$  is known as an  $\ell_2$  norm.

```
library(glmnet)

#initialize matrices of training/test data sets
train.mat=model.matrix(Apps~.,data=train.set) #x-matrix
```

```
test.mat=model.matrix(Apps~.,data=test.set) #y-vector
#initialize grid of possible lambda values, from 10^10 to 10^-2
lambda=10^seq(10,-2,length=100)
#fit ridge regression model (alpha=0 -> ridge regression)
ridge=glmnet(train.mat,train.set$Apps,alpha=0,lambda=lambda)
#perform cross-validation to find best lambda
ridge.cv=cv.glmnet(train.mat,train.set$Apps,alpha=0,lambda=lambda)
lambda.best=ridge.cv$lambda.min
lambda.best #And the best lambda is...
```

```
## [1] 0.01
```

```
#Compute test error
ridge.pred=predict(ridge,test.mat,alpha=0,s=lambda.best)
ridge.MSE=mean((test.set$Apps-ridge.pred)^2)
ridge.MSE
```

```
## [1] 1101724
```

Our test error is slightly lower than that of least squares regression.

d.) Fit a lasso model on the training set, with the tuning parameter chosen by cross-validation, and report the test error obtained, along with the number of non-zero coefficient estimates.

- Background on the Lasso:

Ridge regression isn't perfect- unlike best subset selection and forward/backwards stepwise selection, ridge regression includes all  $p$  predictors in the final model. Even though ridge regression shrinks coefficient estimates, it will never go down to 0, thus eliminating a predictor. Is this a problem for prediction accuracy? Not quite... but it does pose problems for interpretability. The lasso attempts to overcome this disadvantage by being able to generate sparse models. Lasso coefficient estimates seek to minimize the quantity:

$$RSS + \lambda \sum_{i=1}^p |\beta_i|$$

Lasso features an  $\ell_1$  norm;  $||\beta_1|| = \sum |\beta_i|$ . Unlike the  $\ell_2$  norm in ridge regression, the  $\ell_1$  norm in lasso is able to shrink coefficients down to zero when  $\lambda$  is large enough.  $\lambda$  is typically chosen through cross-validation.

```
#coding for lasso works the same way as ridge regression
#only difference: set alpha=1
#we use the same train.mat,test.mat,and lambda grid in part (c)

#fit lasso model
lasso=glmnet(train.mat,train.set$Apps,alpha=1,lambda=lambda)
#perform cross-validation to find best lambda
lasso.cv=cv.glmnet(train.mat,train.set$Apps,alpha=1,lambda=lambda)
lambda.best=lasso.cv$lambda.min
lambda.best
```

```
## [1] 24.77076
```

```
#Compute test error
lasso.pred=predict(lasso,test.mat,alpha=1,s=lambda.best)
lasso.MSE=mean((test.set$Apps-lasso.pred)^2)
lasso.MSE
```

```
## [1] 1144529
```

So far, our test MSEs are in the following order: Ridge Regression < Least Squares < Lasso

e.) Fit a PCR model on the training set, with  $M$  (number of components) chosen by cross-validation, and report the test error obtained, along with the value of  $M$  selected by cross-validation.

- Background on Principal Components Regression (PCR):

When you have a large number of predictors that are closely related to each other, principal components regression aims to summarize this set with a smaller number of representative variables. These representative variables are special in that they are somehow able to capture the variability in the data set. Think about it this way... each of the  $n$  observations lives in a  $p$ -dimensional space, but not all dimensions are interesting. In our case, “interest” is measured by how much the data varies in that dimension. PCR finds the most interesting dimensions for us so we’re not looking at all  $p$ -predictors; the dimension of the problem is reduced as a result.


Here is a link to a beautiful visualization of principal components analysis (PCA). I found this incredibly helpful: <http://setosa.io/ev/principal-component-analysis/>

```
library(pls)

##
## Attaching package: 'pls'

## The following object is masked from 'package:stats':
##
##      loadings
#fit PCR model
pcr=pcr(Apps~.,data=train.set,scale=TRUE,validation="CV")
#We use the plot and summary to choose the best M
validationplot(pcr,val.type="MSEP",xlab="M",ylab="Error")
```

## Apps



```
summary(pcr)
```

```
## Data:      X dimension: 600 17
## Y dimension: 600 1
## Fit method: svdpc
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              4045    4035    2155    2160    2012    1685    1648
## adjCV           4045    4037    2153    2161    2069    1675    1644
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          1650    1657    1576    1565    1572    1574    1582
## adjCV       1646    1653    1571    1561    1568    1569    1579
##      14 comps 15 comps 16 comps 17 comps
## CV          1577    1516    1207    1172
## adjCV       1574    1496    1198    1163
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X          31.18   57.70   64.68   70.32   75.80   80.72   84.34
## Apps       1.01   72.29   72.30   73.85   83.83   84.48   84.49
##      8 comps  9 comps 10 comps 11 comps 12 comps 13 comps 14 comps
## X          87.66   90.57   92.96   95.06   96.87   97.96   98.75
## Apps       84.49   86.01   86.21   86.27   86.27   86.30   86.49
```

```
##      15 comps  16 comps  17 comps
## X      99.37    99.85    100.00
## Apps   90.94    93.23    93.53
```

I'm thinking our optimal M is 10!

Before explaining why M=10, please note that I set a cut-off at M=15. An M value close to the total number of predictors is essentially performing least squares; no dimension reduction occurs in that case. With that being said... from the plot, 10 is the point where the error seems to be at its lowest leveling-off point. From the CV results found in the summary, if you go from 5 to 14 components, there's a dip in estimated error at 10.

```
#Compute error with M=10
pcr.pred=predict(pcr,test.set,ncomp=10)
pcr.MSE=mean((test.set$Apps-pcr.pred)^2)
pcr.MSE
```

```
## [1] 1609171
```

Our MSE from PCR is quite higher than our other methods thus far!


**f.) Fit a PLS model on the training set, with M chosen by cross-validation, and report the test error obtained, along with the value of M selected by cross-validation.**

- Background on Partial Least Squares (PLS):

Partial least squares works the same way as PCR almost, but it makes use of the response vector Y to identify new features. Basically, it's the supervised alternative to PCR. PLS and PCR are at equal levels of performance and advantages, so the decision rests on what type of data you're looking at!

```
#fit PCR model
pls=plsr(Apps~.,data=train.set,scale=TRUE,validation="CV")
#We use the plot and summary to choose the best M
validationplot(pls,val.type="MSEP",xlab="M",ylab="Error")
```

## Apps



```
summary(pls)
```

```
## Data:      X dimension: 600 17
## Y dimension: 600 1
## Fit method: kernelppls
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              4045    1957    1553    1505    1413    1229    1189
## adjCV           4045    1954    1546    1500    1395    1206    1179
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          1176    1172    1167    1165    1166    1166    1164
## adjCV        1167    1163    1159    1157    1157    1158    1156
##      14 comps 15 comps 16 comps 17 comps
## CV           1165    1165    1165    1165
## adjCV        1157    1157    1157    1157
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X          26.45   34.74   62.77   65.08   67.25   72.08   76.70
## Apps       77.50   86.92   88.01   91.42   93.18   93.34   93.39
##      8 comps  9 comps 10 comps 11 comps 12 comps 13 comps 14 comps
## X          79.97   82.64   86.73   88.31   90.78   92.44   94.02
## Apps       93.44   93.48   93.50   93.52   93.53   93.53   93.53
```

```
##      15 comps  16 comps  17 comps
## X      95.90    98.18    100.00
## Apps   93.53    93.53    93.53
```

Like in PCR, I choose M to be 10.

```
#Compute error with M=10
pls.pred=predict(pls,test.set,ncomp=10)
pls.MSE=mean((test.set$Apps-pls.pred)^2)
pls.MSE
```

```
## [1] 1100821
```

**g.) Comment on the results obtained. How accurately can we predict the number of college applications received? Is there much difference among the test errors resulting from these five approaches?**

Here are our test MSEs:

- Least squares - 1101792
- Ridge Regression - 1101724
- Lasso - 1144529
- PCR - 1609171
- PLS- 1100821

From least to greatest, here are the performances of each methodology in terms of minimizing error:

PLS < Ridge Regression < Least Squares < Lasso < PCR